

[ORGANISATION NAME]

Artificial Intelligence Policy

Version 1.0 — Adopted [DATE] · Review annually · Owner: [LEADERSHIP BODY]

How to read this. This is a short document, and short by design. Read it with your team, talk about it, sign it, and revisit it next year. Anything in **[square brackets]** is a prompt for you to complete in your own words.

Why this policy exists

[YOU TO WRITE] — One or two paragraphs in your own voice. Who are you, and why does how you use AI matter to the people you serve? The facts are yours; the framing should be too.

Reference draft: AI has already entered how we work — in the writing, the research, the admin, the day-to-day. This document makes our approach visible and accountable, and grounds it in who we are. We are not afraid of new tools. We are careful about what we build with them.

What we believe about AI

These are not rules. They are convictions — things we hold to be true that shape everything that follows.

- 1. Human judgment is irreplaceable.** AI assists people; it does not replace human judgment, relationships, or responsibility. The final decision, and the accountability for it, stays with a person.
- 2. People before efficiency.** We will not use AI in a way that degrades the quality of human attention we owe the people we serve.
- 3. Truth and verification.** We verify what AI produces before we rely on it, and we never present AI-generated work as something it is not.

4. Privacy and confidentiality. Personal and confidential information about the people we serve does not go into public AI tools — an ethical commitment and a requirement of the NZ Privacy Act 2020.

5. Shared accountability. AI use is not a private matter each person navigates alone. It is held together — set by leadership, owned by everyone.

The policy

Scope. This policy applies to all paid staff, volunteers in significant roles, and anyone acting on behalf of **[ORGANISATION]** in an official capacity. It covers AI use in communications, administration, content creation, service delivery, and decision support.

What we welcome

- Drafting and researching communications, documents, and content — where a person reviews and approves the final output before it goes out in our name.
- Administrative tasks: summarising notes, organising information, drafting correspondence — where no confidential personal data is involved.
- Translation and language support.
- Learning, brainstorming, and exploring ideas.

What we do not do

- Enter confidential or personal information about the people we serve into any public AI tool.
- Present AI-generated content as the original, unassisted work of a person.
- Use AI to replace the human care, presence, or judgment that is central to our work.
- Make significant decisions about people based solely on AI-generated analysis.
- Use AI tools that have not been approved for **[ORGANISATION]** use.

What we are honest about

When AI meaningfully helped produce something, we say so. A brief acknowledgement is enough. **The test:** if someone would feel misled to learn AI was used, we should have been more transparent.

Privacy and data

[ORGANISATION] holds personal information under the NZ Privacy Act 2020. Anything typed into a public AI tool leaves our systems and enters the systems of large technology

companies. Personal details and confidential information must never be processed this way.
When in doubt, leave it out.

Who is responsible

[LEADERSHIP BODY] holds collective responsibility for this policy. Every person is personally responsible for their own AI use within it. If you are unsure whether a use is appropriate, ask before you act.

In practice: five questions

For day-to-day decisions, these are enough:

- 1. Does this involve personal or confidential information about anyone we serve? *If yes, do not use AI.*
- 2. Am I comfortable being transparent about AI's role here, if asked? *If not, rethink the use or disclose it clearly.*
- 3. Have I verified the content for accuracy and appropriateness? *If not, do that first.*
- 4. Is this replacing something that deserves human attention? *If yes, choose the person over the tool.*
- 5. Has this tool been approved for our use? *If not, check first.*

Approved AI tools

Approved tools for **[ORGANISATION]** use. Add or remove as your toolset changes.

Tool	Approved uses	Data caution
ChatGPT / Claude	Drafting, research, admin writing	No personal data
Google Gemini	Drafting, summarising, Workspace tasks	No personal data
Microsoft Copilot	Drafting inside our licensed M365 tenant	Follow tenant data settings
Canva AI	Visual content, graphics	Low risk
Grammarly	Editing communications	No sensitive content
DeepL	Translation	Low risk
[Add your tools here]		

"No personal data" means: never enter names, contact details, financial information, health or confidential notes, or anything shared in confidence about any individual.

(Pair this with the one-page Traffic Light Data Guide.)

Adoption

[YOU TO WRITE] — Two sentences on why you are adopting this. What do you actually want to say? Avoid formulas.

Reference draft: We adopt this policy not as a compliance exercise but as an expression of who we are and how we intend to care for the people we serve. We commit to reviewing it each year as AI changes and as we change.

[LEADER ROLE 1]	_____
[LEADER ROLE 2]	_____
Date adopted	_____
Next review (annual)	_____

Prepared with the AI Policy Starter Kit by Nathaniel Baldock — AI Consulting New Zealand & Global (nathanielbaldock.com). Grounded in the NZ Privacy Act 2020 and international AI governance standards (UNESCO, OECD, ISO/IEC 42001:2023, NIST AI RMF). A starting point — adapt it to your context and review it annually.